

Isha Vohra

US Citizen | 571-425-1316 | ivohra01@gmail.com | linkedin.com/in/ishavohra | <https://isha-vohra.github.io/>

EDUCATION

Georgia Institute of Technology | Atlanta, Georgia

August 2020 to December 2024

BS + MS in Computer Science with Machine Learning Concentration | Highest Honors

Research: Studied the Interplay between Tokenization and Optimization in LLMs under Dr. Kartik Goyal

Examined Image-Caption Fidelity metrics and dataset curation under Dr. James Hays, the creator of Microsoft COCO

TECHNICAL SKILLS

Languages: Java, Python, SQL, JavaScript, R, C++, C, TypeScript

Tools and Frameworks: AWS, GCP, Azure, distributed systems, Docker, Trino, Apache Spark, gRPC, Kafka, Terraform, Databricks, PyTorch, Dynatrace, ELK Stack, Figma, Git, React, Bash, Unix, D3, K8s, AWS S3, Lambda, Redis, Milvus, Node.js

RELEVANT EXPERIENCE

Salesforce | Software Engineer | San Francisco, CA

January 2025 - Present

- Founding engineer on Data360's Data Governance platform: a Postgres-based, tenant-sharded system enforcing metadata and data security across >100M requests/day.
- Built a data observability layer from scratch, using autonomous agents that parse custom log schemas to automatically diagnose root cause and recommend fixes, feeding directly into the on-call paging system to improve platform reliability.
- Fixed a critical (Sev-1) authorization vulnerability by re-architecting Row and Field-Level Security propagation to use eventually-consistent sync, closing a data exposure gap without added latency.
- Root-caused and fixed a 7x latency regression in query processing, restoring baseline throughput at scale.
- Refactored the Functional Integration Test suite (150+ tests) via parallelization and data prehydration, cutting failure rates from 5.6% to 0.03% (99.5% reduction) and average runtime from 104s to 5.6s (94.6% faster).

Salesforce | Software Engineering Intern | San Francisco, CA

May 2024 - Aug. 2024

- Developed a 94.7% accurate ML approach using Microsoft Presidio (spaCy-based NER) to mask personal information from unstructured data.
- Built an Apache Spark pipeline for query validation and translation, eliminating manual intervention.
- Architected cross-team solution involving integration of previously disconnected microservices for authorization, query processing, data ingestion, and dynamic policy-based masking for Attribute Based Access Control (ABAC).

Tableau | Client Platform Engineering Intern | Atlanta, GA

May 2023 - August 2023

- Engineered a data pipeline converting REST API JSON feeds into a production Tableau analytics suite, serving 45K enterprise stakeholders and leading to a 4x increase in efficiency.
- Placed 3rd out of 813 in the enterprise AI Pitch Competition; pitched prototype architecture directly to the CIO.

Northwestern Mutual | Software Engineering Intern | Milwaukee, WI

June 2022 - August 2022

- Designed a dashboard to monitor CPU usage and Docker services/hosts from a state of 0% to 96% coverage.
- Created synthetic monitors to automate testing of key services for the organization using Dynatrace.

PROJECTS

Curio - Webapp | JavaScript, Firebase (Auth, Functions, Firestore), Gemini API, Capacitor

December 2025

- Architected an AI-Driven Gamified Micro-Learning Platform using a Wikipedia random-walk discovery search and the Gemini API to dynamically summarize raw topics into interactive, 1-minute lessons.
- Optimized system reliability and speed by designing a multi-tier cache across Firestore to minimize cloud function invocation costs and latency.

Predicting Traffic Stops and Socio-Economic Impact | Python, Pandas, Folium

December 2024

- Analyzed socio-economic factors affecting US traffic stops to inform budget allocation and identify biases.
- Processed 1M+ traffic stop records by integrating multi-source data with choropleth visualization.
- Achieved 85% accuracy in predicting traffic stop frequencies through regression and neural network models.